Rohail Taimour

Python/R Software Engineer | Data Scientist | Statistician | AI/Machine Learning Specialist

1040, Brussels, Belgium
rohail.taimour@gmail.com | +32 489 83 64 76 | Personal website | Linkedin | Github

Education

MSc in Statistics - KU Leuven, Leuven, Belgium - 2014-2016

Graduated Cum Laude, Master's thesis on continuous optimization of production processes in MATLAB

BSc (Hons.) in Accounting and Finance - Lahore University of Management Sciences (LUMS), Lahore, Pakistan - 2010-2014

Graduated with Distinction (3.6/4.0)
Courses: Operations Research, Supply Chain, Decision analysis, Applied Probability
Treasurer for University Adventure Society organizing hiking trips for groups of upto 300 people in North of Pakistan

Technical competencies

Programming Languages:	Python and R with 7+ years of experience
Cloud Services:	AWS (S3, ECS, SageMaker Studio), Azure (Blob, Databricks, Pipelines)
Data Science/Machine Learning:	PyTorch, Pandas, PyMC3, scikit-Learn, MLflow and standard stack
Data Engineering:	Kedro, Prefect, PySpark
Development Environments:	Pycharm, VScode, Rstudio, Jupyter Notebooks, Azure Databricks
Package Management:	Conda, Mamba, Pip, Poetry for Python package and environment management

Summary

As a seasoned Python/R Software Engineer and Data Scientist with a Master's degree in Statistics, I specialize in creating robust data products with a focus on machine learning. My dynamic skill set bridges the gap between data science, machine learning engineering, and data engineering, allowing me to thrive in fast-paced, ambiguous environments. I am passionate about enhancing project efficiency and impact through best practices and agile methodologies.

Freelance projects (Oct 2022-present)

Automated SQL Script Generation for Cross-Platform Data Migration in PostgreSQL

Python Software Engineer, Illumina, Mechelen, Belgium July-Aug 2023

- Designed and implemented an ORM Mapper for dynamic ingestion of various file formats, automating the SQL script generation process for data migration.
- Implemented the solution as a Python package encapsulating the entire data migration logic within a Docker entrypoint for portability and ease of deployment.
- Conducted comprehensive testing of generated SQL scripts using mock PostgreSQL database tables, ensuring script accuracy and reliability.
- Parameterized key inputs allowing for seamless deployment across multiple environments (Development, Integration, Production).

Multi-Layered Python Solution for Bioinformatics Pipeline Management and Automation

Python Software Engineer and Data Pipeline Architect, Illumina, Mechelen, Belgium April 2023 - October 2023

- Implemented a Python service to routinely monitor new sequencing data, tracking progress using an SQLite database.
- Developed a multi-layered Python package: one layer encapsulated a data processing component, packaged within a **Docker** runtime environment, for handling bioinformatics pipeline outputs.
- Created a separate entry point within the package responsible for initiating and managing the Docker-packaged data processing component, as well as other routine operations such as downloading completed analyses, updating analyses statuses, etc
- Implemented comprehensive systems integration, utilizing a combination of CLI tools and API calls for effective coordination and automation across various software components.
- Applied Object-Oriented Programming (OOP) techniques to organize API, database interactions and endpoint processing.
- Implemented unit testing using **pytest** and implemented fail-safe mechanisms for robust error handling.

Design and implement information retrieval methods using Natural language processing (NLP)

Machine Learning Engineer, IT Supply Quality, GSK Belgium Oct 2022-Feb 2023

- Improved performance of information retrieval by 20% on unseen test data using a custom named entity recognition (NER) from **Spacy**.
- Performed POC's on Azure DataBricks environment to improve model performance using rule-based techniques as well as **NER** and annotated data to train custom NER.
- Added text preprocessing features to the NLP pipeline such as **Spacy** tokenization, Part of speech (POS) tagging, better handling of non-english emails, breaking emails into sentences, etc.

CI/CD:	Git, GitHub Actions, Azure Pipelines, GitLab Pipelines, GitHub CLI
Containerization:	DockerHub, Docker
Database Management:	PostgreSQL, SQLite3, Neo4j, SQLAlchemy
Technical Documentation:	Pandoc, Markdown, Sphinx for documentation; CSS, HTML for web development
Software Testing:	Pytest for testing; Black, Pre-Commit, iSort, Flake8, Mypy for code quality

Personal details

- Nationality: Belgian, Pakistani
- Languages: English (fluent/bilingual), Urdu (Native), French (B1)
- Mobility: Driving License available, flexible for hybrid setup in Belgium
- Availability: Immediately
- Hobbies: Drumming and percussion instruments, Bouldering/Climbing, productivity, Squash, reading

Data science projects at IT AI team, UCB Pharmaceutical (2016-Oct 2022)

Yield optimization for batch and continuous production processes using Machine Learning in Python

Lead Data Scientist, Supply and Manufacturing, UCB Switzerland/Belgium Aug 2020-Oct 2022

- Production setting proposed by model directly led to an increased throughput of 20%, turning in a recurring 1.5 million euro in annual cost savings
- Analyze time series data collected from equipment sensors and visually summarize golden batch insights
- Created (Bayesian) and tree-based regression models to quantify impact of process changes and predict batch performance
- Performed a thorough model validation and hyperparameter tuning exercise before recommending model insights be tested in a live production environment
- Supported delivery of workshops demystifying the process of conducting AI projects and machine learning to process experts

Python Framework for Customized Promotional Responsiveness Models Across Regions

AI/ML engineer, Lead Data Scientist, Go to Market/Commercial EU5, US and Japan, UCB June 2019-June 2021

- Developed a Python package with **Cookiecutter** templates that abstract the complexities of the data science workflow, enabling configurable deployments across diverse scenarios such as different countries and disease areas.
- Enhanced the package to seamlessly wrap over **scikit-learn**, thereby simplifying key data science tasks from preprocessing to model training and tuning
- Incorporated **MLflow** into the package for robust artifact management, allowing for the tracking of model versions, data inputs, and predictions
- Created customer segmentation models and proposed optimal resource allocation based on customer responsiveness to different marketing channels
- Investigated adaptations to data science methodology for country/product specificities for maximum reusability. Delivered as many as ten different use cases for different products and countries
- Performed feature engineering using **PySpark** and validated ingested data using data visualization methods and discussions with subject-matter experts

Scientific influencer (KOLs) identification, ranking and profiling using network analytics and Neo4j

Data scientist/Product owner, Drug Development, Commercial, Medical affairs, UCB 2018-2019

- Developed custom **Neo4j** databases integrating diverse data sources for KOL influence analysis, enhancing data-driven decision-making.
- Utilized **py2neo** within **Jupyter Notebooks** for interactive data manipulation and network visualizations, employing tools like NetworkX and **Cytoscape** for insightful analysis.
- Interacted with the Graph database via **Cypher** queries in the web UI as well as via the CLI for data extraction, exploration and reporting.
- Supported improvements in the intake of customer requests to reduce time to deliver reports from days to hours

Developed an Automated Forecasting Workflow of Claims Data from US Healthcare System

Lead Data Scientist, US Finance and claims, UCB 2017-2018

- Engineered specialized **R packages** focusing on separate concerns: data engineering for preprocessing, a wrapper over Facebook's Prophet for advanced forecasting, and automated

reporting for performance analysis.
- Designed and implemented a comprehensive end-to-end workflow for ingesting healthcare claims data, performing time-series forecasting, and generating insightful reports on forecasting accuracy.
- Achieved over 90% forecasting accuracy across various use cases by meticulously tuning models and integrating bespoke anomaly detection algorithms for time series data.
- Conducted extensive hyperparameter tuning and model validation using **high-performance computing** to optimize forecasting models effectively.
- Automated report generation using **R Markdown**, providing clear, concise insights into forecasting performance and model accuracy.

Hands-on workshop to demystify Artificial intelligence and Machine Learning

Data science instructor, IT departments US, EU, UCB May- June 2017

- Created a R shiny application to create an engaging way for participants to learn about typical AI use cases
- Delivered the workshop to over 100 people in four different venues and received great feedback on level of engagement

Personal projects

Web Scraper to analyse Property Purchase and Rental Trends in Belgium

- Developed web scraper using Beautiful Soup to collect information such as apartment data such as price, area, etc.
- Implemented SQLite for data storage, using `**Pydantic**` for data validation and `**SQLAlchemy**` for database interactions.
- Encapsulated the concerns into a python package with dependency management using Poetry.
- Employed **Prefect** for job orchestration, managing the workflow's scheduling and monitoring of scraping tasks.

Personal Portfolio and blogging website built using Hugo and hosted using Github Pages

- Created website using **Hugo** and implemented features such as a contact form, and visitor commenting capabilities.
- Hosted the static website on GitHub Pages and automated the deployment process using GitHub Actions.
- Codebase hosted on github

Automated Resume Builder and Continuous Deployment System with GitHub Pages Hosting

- Engineered an automated system for generating, versioning, and hosting a dynamic CV using Markdown, HTML, Jinja templating and CSS.
- Set up a trio of GitHub repositories to separately manage the CV's content, styling, and public hosting on Github Pages.
- Developed a Python package for automating the styling and generation of the CV, integrating with Markdown and HTML/CSS.
- Implemented version control for CV content using a private GitHub repository, ensuring secure and organized data management.
- Leveraged GitHub Actions for automating the CV's generation and deployment process, enabling updates through git pushes.
- Hosted the final CV on GitHub Pages, providing a live, online version that can be easily updated

Unit Commitment Solver for Power Grid Optimization via FastAPI

- Developed a REST API using **FastAPI** for optimizing energy distribution among powerplants based on load requirements and fuel costs.
- Implemented multiple algorithms to solve the **unit-commitment problem**, considering factors like fuel cost, powerplant efficiency, and environmental constraints.
- Utilized **Pydantic** for data validation and schema definition, ensuring data integrity and streamlined request handling.

- Packaged and containerized the application using **Docker**, with detailed documentation and a Dockerfile for easy deployment and scalability.
- Employed **pytest**, along with Python best practices such as typing and linting.
- Managed project dependencies using **Poetry**, facilitating efficient workflow and package management.
- Deployed the API service using **Uvicorn** and integrated a **Swagger UI** for interactive API documentation and testing